

Autoscaling LAB (HPA)

03 January 2021 16:05

- Setup the lab environment, in my case used the minikube

```
$ minikube status
minikube
type: Control Plane
host: Running
kubelet: Running
apiserver: Running
kubeconfig: Configured
timeToStop: Nonexistent

$ minikube addons enable metrics-server
```

- Create a deployment as follow

```
#cat testdep.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: php-apache
spec:
  selector:
    matchLabels:
      app: php-apache
  replicas: 1
  template:
    metadata:
      labels:
        app: php-apache
    spec:
      containers:
        - name: php-apache
          image: gcr.io/google_containers/hpa-example
      resources:
        requests:
          cpu: "500m"
          memory: "500M"
        ports:
          - containerPort: 80
```

- Expose the deployment

```
$ kubectl expose deployment php-apache --name=php-apache --port=80
service/php-apache exposed
```

- Test the deployment is serving using another test pod

```
$ kubectl run -i --tty service-test --image=busybox /bin/sh
wget -q -O- http://php-apache.default.svc.cluster.local
OK!/#

kubectl exec -i --tty service-test /bin/sh
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
/# wget -q -O- http://php-apache.default.svc.cluster.local
OK!/#
```

- Define the HPA (Horizontal pod autoscale)

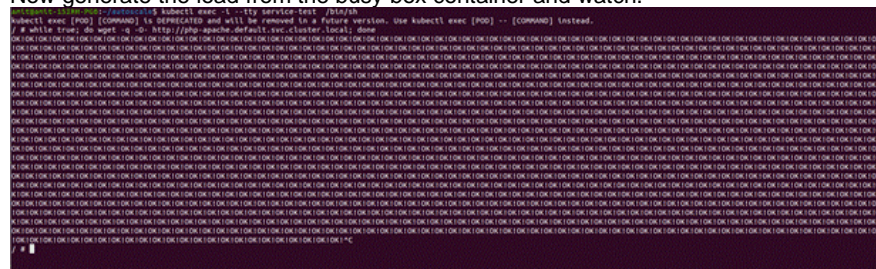
```
$ kubectl get hpa
No resources found in default namespace.

$ kubectl autoscale deployment php-apache --cpu-percent=50 --min=1 --max=3
horizontalpodautoscaler.autoscaling/php-apache autoscaled

$ kubectl get hpa
NAME           REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
php-apache     Deployment/php-apache  0/50%    1         3         0         1s
```

- sdfs
- Now generate the load from the busy box container and watch.

```
kubectl exec -i --tty service-test /bin/sh
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
# while true; do wget -q -O- http://php-apache.default.svc.cluster.local; done
```



Parallely on the second terminal observe pods and HPA output. After 2minutes you will notice more pods being launched

```

amit@amit-15IRH-PG0:~/autoscale$ kubectl top pods
NAME          CPU(cores)   MEMORY(bytes)
php-apache-59f68b8bcc-v5dcp  735m        12Mi
service-test  8m          1Mi
amit@amit-15IRH-PG0:~/autoscale$ kubectl top pods
NAME          CPU(cores)   MEMORY(bytes)
php-apache-59f68b8bcc-v5dcp  735m        12Mi
service-test  8m          1Mi
amit@amit-15IRH-PG0:~/autoscale$ kubectl get hpa
NAME          REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
php-apache    Deployment/php-apache  147%/50%   1         3         3         11m
amit@amit-15IRH-PG0:~/autoscale$

```

```

amit@amit-15IRH-PG0:~/autoscale$ kubectl get all
NAME          READY   STATUS    RESTARTS   AGE
pod/php-apache-59f68b8bcc-9wzbs  1/1     Running   0          5m25s
pod/php-apache-59f68b8bcc-jn2mv  0/1     Pending   0          5m25s
pod/php-apache-59f68b8bcc-v5dcp  1/1     Running   0          29m
pod/service-test                 1/1     Running   1          27m

NAME          TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
service/kubernetes  ClusterIP    10.96.0.1    <none>        443/TCP    75m
service/php-apache  ClusterIP    10.98.84.229 <none>        80/TCP     18m

NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/php-apache  2/3     3            2          29m

NAME          DESIRED   CURRENT   READY   AGE
replicaset.apps/php-apache-59f68b8bcc  3         3         2       29m

NAME          REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
horizontalpodautoscaler.autoscaling/php-apache  Deployment/php-apache  0%/50%   1         3         3         16m
amit@amit-15IRH-PG0:~/autoscale$

```

Now press ctrl + C to stop more load generation in busy box container and after few minutes you will notice additional pods will disappear

```

amit@amit-15IRH-PG0:~/autoscale$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
php-apache-59f68b8bcc-v5dcp  1/1     Running   0          30m
service-test  1/1     Running   1          29m
amit@amit-15IRH-PG0:~/autoscale$ kubectl get hpa
NAME          REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
php-apache    Deployment/php-apache  0%/50%   1         3         3         17m
amit@amit-15IRH-PG0:~/autoscale$ kubectl get hpa
NAME          REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
php-apache    Deployment/php-apache  0%/50%   1         3         1         17m
amit@amit-15IRH-PG0:~/autoscale$

```