

Installation LAB

16 December 2020 17:12

ISO URL - http://centos.excellmedia.net/7.9.2009/isos/x86_64/CentOS-7-x86_64-Minimal-2009.iso

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

• Pre-requisite

- 2 GB or more of RAM per machine (any less will leave little room for your apps)
- 2 CPUs or more
- Full network connectivity between all machines in the cluster (public or private network is fine)
- Unique hostname, MAC address, and product_uuid for every node. See [here](#) for more details.
 - You can get the MAC address of the network interfaces using the command `ip link` or `ifconfig -a`
 - The product_uuid can be checked by using the command `sudo cat /sys/class/dmi/id/product_uuid` (VM)
- Certain ports are open on your machines. See [here](#) for more details.
- Swap disabled. You **MUST** disable swap in order for the kubelet to work properly.

From <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

As a requirement for your Linux Node's iptables to correctly see bridged traffic, you should ensure `net.bridge.bridge-nf-call-iptables` is set to 1 in your `sysctl` config, e.g.

ON ALL YOUR MACHINES (VMs)

Prepare the VMs (Master & Worker):

There are a few things you need to do to get the VM ready. Specifically, you need to turn off swap, tweak some configuration settings, and make sure you have the prerequisites libraries installed. To do that, follow these steps:

1. Networking (Assign IP address [static])

```
# nmcli -p dev # List network devices
# nmtui edit enps03
- 192.168.1.76 g/w 192.168.1.1
```

2. Change to root:

```
sudo su
```

3. Set hostname and add entries into /etc/hosts

a. Master:

```
# hostnamectl set-hostname k8s-master
[root@k8s-master ~]# cat /etc/hosts
192.168.1.73 k8s-master
192.168.1.74 worker1
```

b. Workers

```
# hostnamectl set-hostname worker1
# hostnamectl set-hostname worker2
[root@worker1 ~]# cat /etc/hosts
192.168.1.73 k8s-master
192.168.1.74 worker1
```

4. Turn off swap: To do this, you will first need to turn it off directly ...

```
swapoff -a
```

... then comment out the reference to swap in /etc/fstab. Start by editing the file:

a. `vi /etc/fstab`

5. Letting iptables see bridged traffic

- Make sure that the `br_netfilter` module is loaded.

```
[root@localhost ~]# lsmod | grep br_netfilter
[root@localhost ~]# modprobe br_netfilter
[root@localhost ~]# lsmod | grep br_netfilter
br_netfilter 22256 0
bridge 151336 2 br_netfilter,ebtable_bridge
```

- Disable `firewalld` [optional - else we need to open specific ports]

```
[root@k8s-master ~]# systemctl stop firewalld
[root@k8s-master ~]# systemctl disable firewalld
Removed symlink /etc/systemd/system/multi-user.target.wants/firewalld.service.
Removed symlink /etc/systemd/system/dbus-org.fedoraproject.FirewallD1.service.
```

- As a requirement for your Linux Node's iptables to correctly see bridged traffic, you should ensure `net.bridge.bridge-nf-call-iptables` is set to 1 in your `sysctl`

```
cat <<EOF | sudo tee net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-iptables = 1
/etc/sysctl.d/k8s.conf
EOF
sudo sysctl --system
```

6. Reboot so the changes take effect [optional]

Installing runtime (docker/CRI-O/Containerd)

From <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

```
# (Install Docker CE)
## Set up the repository
```

Install required packages

```
sudo yum install -y yum-utils device-mapper-persistent-data lvm2
```

Add the Docker repository

```
sudo yum-config-manager --add-repo \
https://download.docker.com/linux/centos/docker-ce.repo
```

Install Docker CE

```
sudo yum update -y && sudo yum install -y \
containerd.io-1.2.13 \
docker-ce-19.03.11 \
docker-ce-cli-19.03.11
```

Create /etc/docker

```
sudo mkdir /etc/docker
```

Set up the Docker daemon

```
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2",
  "storage-opts": [
    "overlay2.override_kernel_check=true"
  ]
}
EOF
```

Create /etc/systemd/system/docker.service.d

```
sudo mkdir -p /etc/systemd/system/docker.service.d
```

Restart Docker

```
sudo systemctl daemon-reload
sudo systemctl restart docker
If you want the docker service to start on boot, run the following command:
sudo systemctl enable docker
```

Installing kubeadm, kubelet and kubectl

- Add repository

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-$basearch
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
exclude=kubelet kubeadm kubectl
EOF
```

- Set SELinux to permissive

```
sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

- Start kubelet service

```
sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
sudo systemctl enable --now kubelet
```

Master Node

```
[root@k8s_master ~]# docker image ls
REPOSITORY TAG IMAGE ID CREATED SIZE
[root@k8s_master ~]# kubeadm config images pull
```

```
[root@k8s_master ~]# kubeadm config images pull >/dev/null 2>&1
[root@k8s_master ~]# docker image ls
REPOSITORY TAG IMAGE ID CREATED SIZE
k8s.gcr.io/kube-proxy v1.20.1 e3f6fcd87756 6 days ago 118MB
k8s.gcr.io/kube-apiserver v1.20.1 75c7f7112080 6 days ago 122MB
k8s.gcr.io/kube-controller-manager v1.20.1 2893d78e47dc 6 days ago 116MB
k8s.gcr.io/kube-scheduler v1.20.1 4aa0b4397bbb 6 days ago 46.4MB
k8s.gcr.io/etcd 3.4.13-0 0369cf4303ff 3 months ago 253MB
k8s.gcr.io/coredns 1.7.0 bfe3a36ebd25 6 months ago 45.2MB
k8s.gcr.io/pause 3.2 80d28bedfe5d 10 months ago 683kB
```

--apiserver-advertise-address	The IP address the API Server will advertise it's listening on. If not set the default network interface will be used.
--image-repository	Choose a container registry to pull control plane images from (default "k8s.gcr.io")
--kubernetes-version	Choose a specific Kubernetes version for the control plane. (default "stable-1")
--pod-network-cidr	Specify range of IP addresses for the pod network. If set, the control plane will automatically allocate CIDRs for every node.
--service-cidr	Use alternative range of IP address for service VIPs. (default "10.96.0.0/12")

```
# kubeadm init --apiserver-advertise-address=<Master_IP> --pod-network-cidr=192.168.0.0/16
```

Master IP = Internal IP Address

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 192.168.1.73:6443 --token clrxr9.qelzztg0zns24hwu \
--discovery-token-ca-cert-hash sha256:31a30fdb81a7ff3bcd0c17234b0da37b10b700c5fcfa0e3f2bf2f8e478322e85
```

Note: Currently Calico is the only CNI plugin that the kubeadm project performs e2e tests against. If you find an issue related to a CNI plugin you should log a ticket in its respective issue tracker instead of the kubeadm or kubernetes issue trackers.

From <<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/#pod-network>>

POD Network

```
[root@k8s-master ~]# kubectl get pods -A
NAMESPACE      NAME                                READY   STATUS    RESTARTS   AGE
kube-system    coredns-74ff55c5b-2dfzq            0/1     Pending   0           10m
kube-system    coredns-74ff55c5b-h4nw5            0/1     Pending   0           10m
kube-system    etcd-k8s-master                    1/1     Running   0           10m
kube-system    kube-apiserver-k8s-master           1/1     Running   0           10m
kube-system    kube-controller-manager-k8s-master 1/1     Running   0           10m
kube-system    kube-proxy-pn8rc                    1/1     Running   0           10m
kube-system    kube-scheduler-k8s-master          1/1     Running   0           10m
[root@k8s-master ~]#
```

```
[root@k8s-master ~]# kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
podsecuritypolicy.policy/psp.flannel.unprivileged created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
```

```
[root@k8s-master ~]# kubectl get pods -A
NAMESPACE      NAME                                READY   STATUS    RESTARTS   AGE
kube-system    coredns-74ff55c5b-2dfzq            0/1     Running   0           11m
kube-system    coredns-74ff55c5b-h4nw5            1/1     Running   0           11m
kube-system    etcd-k8s-master                    1/1     Running   0           12m
kube-system    kube-apiserver-k8s-master           1/1     Running   0           12m
kube-system    kube-controller-manager-k8s-master 1/1     Running   0           12m
kube-system    kube-flannel-ds-tksk6               1/1     Running   0           44s
kube-system    kube-proxy-pn8rc                    1/1     Running   0           11m
kube-system    kube-scheduler-k8s-master          1/1     Running   0           12m
[root@k8s-master ~]#
```

Worker Nodes

```
-----
[root@worker1 ~]# kubeadm join 192.168.1.73:6443 --token clrxr9.qelzztg0zns24hwu \
> --discovery-token-ca-cert-hash sha256:31a30fdb81a7ff3bcd0c17234b0da37b10b700c5fcfa0e3f2bf2f8e478322e85
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...
```

This node has joined the cluster:

- * Certificate signing request was sent to apiserver and a response was received.
- * The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

```
[root@worker1 ~]#
```

```
[root@k8s-master ~]# kubectl get nodes
NAME        STATUS    ROLES    AGE   VERSION
k8s-master  Ready     control-plane,master   14m   v1.20.1
worker1     NotReady  <none>    48s   v1.20.1
[root@k8s-master ~]#
```