

Demo

- If you were deploying on a VMware environment NSX-T may be a good option.
- If you look at the play-with-k8s labs they use WeaveNet.
- In our demos in the course we used Calico.
- Depending on your environment and after evaluating the Pros and Cons of each of these, you may chose the right networking sol ution.

Reset your cluster:

```
[root@k8s-master ~]# kubeadm init --apiserver-advertise-address=192.168.1.73 --pod-network-cidr=192.168.0.0/16
[init] Using Kubernetes version: v1.20.1
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
[certs] Using certificateDir folder '/etc/kubernetes/pki'
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [k8s-master kubernet.es kubernet.es.default kubernet.es.default.svc kubernet.es.default.svc.cluster.local] and IPs
[10.96.0.1 192.168.1.73]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
```

Initialize Cluster

```
[root@k8s-master ~]# kubectl get pods -A
NAMESPACE      NAME                                     READY   STATUS    RESTARTS   AGE
kube-system     coredns-74ff55c5b-bt6lb                0/1     Pending   0           3m24s
kube-system     coredns-74ff55c5b-xqkt7                0/1     Pending   0           3m24s
kube-system     etcd-k8s-master                        1/1     Running   0           3m37s
kube-system     kube-apiserver-k8s-master              1/1     Running   0           3m37s
kube-system     kube-controller-manager-k8s-master     1/1     Running   0           3m37s
kube-system     kube-proxy-v8wvk                      1/1     Running   0           3m24s
kube-system     kube-scheduler-k8s-master             1/1     Running   0           3m37s
[root@k8s-master ~]#
```

Join workers without installing CNI and notice status:

```
[root@k8s-master ~]# kubectl get pods -A -w
NAMESPACE      NAME                                     READY   STATUS    RESTARTS   AGE
kube-system     coredns-74ff55c5b-bt6lb                0/1     Pending   0           11m
kube-system     coredns-74ff55c5b-xqkt7                0/1     Pending   0           11m
kube-system     etcd-k8s-master                        1/1     Running   0           11m
kube-system     kube-apiserver-k8s-master              1/1     Running   0           11m
kube-system     kube-controller-manager-k8s-master     1/1     Running   0           11m
kube-system     kube-proxy-d2rcc                      1/1     Running   0           4m44s
kube-system     kube-proxy-mkz62                      1/1     Running   0           4m55s
kube-system     kube-proxy-v8wvk                      1/1     Running   0           11m
kube-system     kube-scheduler-k8s-master             1/1     Running   0           11m
^C[root@k8s-master ~]# kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
k8s-master          NotReady control-plane,master 11m   v1.20.1
worker26            NotReady <none>    4m59s v1.20.1
worker6             NotReady <none>    4m48s v1.20.1
[root@k8s-master ~]#
```

Install network drivers**FLANNEL**

```
kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

Kubernetes Scheduling:

```
[root@k8s-master ~]# kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
k8s-master          Ready     control-plane,master 29m   v1.20.1
worker6             Ready     <none>    22m   v1.20.1
[root@k8s-master ~]# kubectl get pods -A -o wide
NAMESPACE      NAME                                     READY   STATUS    RESTARTS   AGE   IP              NODE      NOMINATED NODE   READINESS GATES
default        net-test-dc8bf8666-2g4js              1/1     Running   0           10m   192.168.2.8     worker6    <none>            <none>
default        net-test-dc8bf8666-4plqc              1/1     Running   0           10m   192.168.2.5     worker6    <none>            <none>
default        net-test-dc8bf8666-9rbz4              1/1     Running   0           10m   192.168.2.6     worker6    <none>            <none>
default        net-test-dc8bf8666-bvbf7              1/1     Running   0           10m   192.168.2.4     worker6    <none>            <none>
default        net-test-dc8bf8666-jcdtt              1/1     Running   0           10m   192.168.2.7     worker6    <none>            <none>
default        net-test-dc8bf8666-jmx49              1/1     Running   0           10m   192.168.2.3     worker6    <none>            <none>
kube-system     coredns-74ff55c5b-bt6lb                0/1     ContainerCreating 0           29m   <none>          worker26   <none>            <none>
kube-system     coredns-74ff55c5b-xqkt7                0/1     Running   0           29m   192.168.2.2     worker6    <none>            <none>
kube-system     etcd-k8s-master                        1/1     Running   0           29m   192.168.1.73    k8s-master <none>            <none>
kube-system     kube-apiserver-k8s-master              1/1     Running   0           29m   192.168.1.73    k8s-master <none>            <none>
kube-system     kube-controller-manager-k8s-master     1/1     Running   0           29m   192.168.1.73    k8s-master <none>            <none>
kube-system     kube-flannel-ds-g7tgm                  1/1     Running   1           16m   192.168.1.75    worker26   <none>            <none>
kube-system     kube-flannel-ds-hmjvs                  1/1     Running   0           16m   192.168.1.73    k8s-master <none>            <none>
kube-system     kube-flannel-ds-mjkfg                  1/1     Running   0           16m   192.168.1.89    worker6    <none>            <none>
kube-system     kube-proxy-d2rcc                      1/1     Running   0           22m   192.168.1.89    worker6    <none>            <none>
kube-system     kube-proxy-mkz62                      1/1     Running   1           22m   192.168.1.75    worker26   <none>            <none>
kube-system     kube-proxy-v8wvk                      1/1     Running   0           29m   192.168.1.73    k8s-master <none>            <none>
kube-system     kube-scheduler-k8s-master             1/1     Running   0           29m   192.168.1.73    k8s-master <none>            <none>
[root@k8s-master ~]#
```

```
[root@k8s-master ~]# kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
k8s-master          Ready     control-plane,master 31m   v1.20.1
worker26            Ready     <none>    104s  v1.20.1
worker6             Ready     <none>    24m   v1.20.1
[root@k8s-master ~]#
```

```
[root@k8s-master ~]# kubectl scale deployment --replicas=8 net-test
deployment.apps/net-test scaled
[root@k8s-master ~]# kubectl get pods -A -o wide
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
default	net-test-dc8bf8666-2g4js	1/1	Running	0	13m	192.168.2.8	worker6	<none>	<none>
default	net-test-dc8bf8666-4plqc	1/1	Running	0	13m	192.168.2.5	worker6	<none>	<none>
default	net-test-dc8bf8666-9rbz4	1/1	Running	0	13m	192.168.2.6	worker6	<none>	<none>
default	net-test-dc8bf8666-bvbf7	1/1	Running	0	13m	192.168.2.4	worker6	<none>	<none>
default	net-test-dc8bf8666-cq22t	0/1	ContainerCreating	0	3s	<none>	worker26	<none>	<none>
default	net-test-dc8bf8666-fngls	0/1	ContainerCreating	0	3s	<none>	worker26	<none>	<none>
default	net-test-dc8bf8666-jcdtt	1/1	Running	0	13m	192.168.2.7	worker6	<none>	<none>
default	net-test-dc8bf8666-pmx49	1/1	Running	0	13m	192.168.2.3	worker6	<none>	<none>
kube-system	coredns-74ff55c5b-bt6lb	0/1	Running	0	32m	192.168.3.2	worker26	<none>	<none>
kube-system	coredns-74ff55c5b-xqkt7	0/1	Running	0	32m	192.168.2.2	worker6	<none>	<none>
kube-system	etcd-k8s-master	1/1	Running	0	32m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-apiserver-k8s-master	1/1	Running	0	32m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-controller-manager-k8s-master	1/1	Running	0	32m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-flannel-ds-g7tgm	1/1	Running	0	18m	192.168.1.75	worker26	<none>	<none>
kube-system	kube-flannel-ds-hmjvs	1/1	Running	0	18m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-flannel-ds-mjkfg	1/1	Running	0	18m	192.168.1.89	worker6	<none>	<none>
kube-system	kube-proxy-d2rcc	1/1	Running	0	25m	192.168.1.89	worker6	<none>	<none>
kube-system	kube-proxy-mkz62	1/1	Running	0	25m	192.168.1.75	worker26	<none>	<none>
kube-system	kube-proxy-v8wvk	1/1	Running	0	32m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-scheduler-k8s-master	1/1	Running	0	32m	192.168.1.73	k8s-master	<none>	<none>

```
[root@k8s-master ~]#
```

Change the Network: CALICO

- Delete flannel

```
kubectl delete -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

- Check status of pods

```
[root@k8s-master ~]# kubectl delete -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

```
podsecuritypolicy.policy "psp.flannel.unprivileged" deleted
clusterrole.rbac.authorization.k8s.io "flannel" deleted
clusterrolebinding.rbac.authorization.k8s.io "flannel" deleted
serviceaccount "flannel" deleted
configmap "kube-flannel-cfg" deleted
daemonset.apps "kube-flannel-ds" deleted
```

```
[root@k8s-master ~]# kubectl get pods -A -o wide
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
default	net-test-dc8bf8666-2g4js	1/1	Running	0	14m	192.168.2.8	worker6	<none>	<none>
default	net-test-dc8bf8666-4plqc	1/1	Running	0	14m	192.168.2.5	worker6	<none>	<none>
default	net-test-dc8bf8666-9rbz4	1/1	Running	0	14m	192.168.2.6	worker6	<none>	<none>
default	net-test-dc8bf8666-bvbf7	1/1	Running	0	14m	192.168.2.4	worker6	<none>	<none>
default	net-test-dc8bf8666-cq22t	1/1	Running	0	86s	192.168.3.4	worker26	<none>	<none>
default	net-test-dc8bf8666-fngls	1/1	Running	0	86s	192.168.3.3	worker26	<none>	<none>
default	net-test-dc8bf8666-jcdtt	1/1	Running	0	14m	192.168.2.7	worker6	<none>	<none>
default	net-test-dc8bf8666-pmx49	1/1	Running	0	14m	192.168.2.3	worker6	<none>	<none>
kube-system	coredns-74ff55c5b-bt6lb	0/1	Running	0	33m	192.168.3.2	worker26	<none>	<none>
kube-system	coredns-74ff55c5b-xqkt7	0/1	Running	0	33m	192.168.2.2	worker6	<none>	<none>
kube-system	etcd-k8s-master	1/1	Running	0	33m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-apiserver-k8s-master	1/1	Running	0	33m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-controller-manager-k8s-master	1/1	Running	0	33m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-flannel-ds-hmjvs	0/1	Terminating	0	20m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-flannel-ds-mjkfg	0/1	Terminating	0	20m	192.168.1.89	worker6	<none>	<none>
kube-system	kube-proxy-d2rcc	1/1	Running	0	26m	192.168.1.89	worker6	<none>	<none>
kube-system	kube-proxy-mkz62	1/1	Running	0	27m	192.168.1.75	worker26	<none>	<none>
kube-system	kube-proxy-v8wvk	1/1	Running	0	33m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-scheduler-k8s-master	1/1	Running	0	33m	192.168.1.73	k8s-master	<none>	<none>

```
[root@k8s-master ~]#
```

- Install Calico

```
[root@k8s-master ~]# ls calico.yaml
```

```
calico.yaml
```

```
[root@k8s-master ~]# kubectl create -f calico.yaml
```

```
configmap/calico-config created
```

```
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
```

```
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
```

```
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
```

```
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
```

```
clusterrole.rbac.authorization.k8s.io/calico-node created
```

```
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
```

```
daemonset.apps/calico-node created
```

```
serviceaccount/calico-node created
```

```
deployment.apps/calico-kube-controllers created
```

```
serviceaccount/calico-kube-controllers created
```

```
poddistributionpolicy.policy/calico-kube-controllers created
```

```
[root@k8s-master ~]#
```

- Check status of Pods

```
[root@k8s-master ~]# kubectl get pods -A -o wide
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
default	net-test-dc8bf8666-2g4js	1/1	Running	0	16m	192.168.2.8	worker6	<none>	<none>
default	net-test-dc8bf8666-4plqc	1/1	Running	0	16m	192.168.2.5	worker6	<none>	<none>
default	net-test-dc8bf8666-9rbz4	1/1	Running	0	16m	192.168.2.6	worker6	<none>	<none>
default	net-test-dc8bf8666-bvbf7	1/1	Running	0	16m	192.168.2.4	worker6	<none>	<none>
default	net-test-dc8bf8666-cq22t	1/1	Running	0	3m21s	192.168.3.4	worker26	<none>	<none>
default	net-test-dc8bf8666-fngls	1/1	Running	0	3m21s	192.168.3.3	worker26	<none>	<none>
default	net-test-dc8bf8666-jcdtt	1/1	Running	0	16m	192.168.2.7	worker6	<none>	<none>
default	net-test-dc8bf8666-pmx49	1/1	Running	0	16m	192.168.2.3	worker6	<none>	<none>
kube-system	calico-kube-controllers-744cdf676-zl22l	0/1	Running	1	45s	192.168.3.5	worker26	<none>	<none>
kube-system	calico-node-bnj6m	1/1	Running	0	45s	192.168.1.73	k8s-master	<none>	<none>
kube-system	calico-node-ns99h	1/1	Running	0	45s	192.168.1.89	worker6	<none>	<none>
kube-system	calico-node-grtbh	0/1	PodInitializing	0	45s	192.168.1.75	worker26	<none>	<none>
kube-system	coredns-74ff55c5b-bt6lb	0/1	Running	0	35m	192.168.3.2	worker26	<none>	<none>
kube-system	coredns-74ff55c5b-xqkt7	0/1	Running	0	35m	192.168.2.2	worker6	<none>	<none>
kube-system	etcd-k8s-master	1/1	Running	0	35m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-apiserver-k8s-master	1/1	Running	0	35m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-controller-manager-k8s-master	1/1	Running	0	35m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-proxy-d2rcc	1/1	Running	0	28m	192.168.1.89	worker6	<none>	<none>
kube-system	kube-proxy-mkz62	1/1	Running	0	29m	192.168.1.75	worker26	<none>	<none>
kube-system	kube-proxy-v8wvk	1/1	Running	0	35m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-scheduler-k8s-master	1/1	Running	0	35m	192.168.1.73	k8s-master	<none>	<none>

```
[root@k8s-master ~]#
```

- Scale up your deployment
- Check IP address of new pods

```
[root@k8s-master ~]# kubectl get pods -A -o wide
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
default	net-test-dc8bf8666-2g4js	1/1	Running	0	18m	192.168.2.8	worker6	<none>	<none>
default	net-test-dc8bf8666-4plqc	1/1	Running	0	18m	192.168.2.5	worker6	<none>	<none>
default	net-test-dc8bf8666-9rbz4	1/1	Running	0	18m	192.168.2.6	worker6	<none>	<none>
default	net-test-dc8bf8666-bvbf	1/1	Running	0	18m	192.168.2.4	worker6	<none>	<none>
default	net-test-dc8bf8666-cq2zt	1/1	Running	0	4m54s	192.168.3.4	worker26	<none>	<none>
default	net-test-dc8bf8666-fngls	1/1	Running	0	4m54s	192.168.3.3	worker26	<none>	<none>
default	net-test-dc8bf8666-jcdtt	1/1	Running	0	18m	192.168.2.7	worker6	<none>	<none>
default	net-test-dc8bf8666-klmct	1/1	Running	0	16s	192.168.187.194	worker26	<none>	<none>
default	net-test-dc8bf8666-kmm9m	1/1	Running	0	16s	192.168.187.193	worker26	<none>	<none>
default	net-test-dc8bf8666-pmx49	1/1	Running	0	18m	192.168.2.3	worker6	<none>	<none>
kube-system	calico-kube-controllers-744cfd676-zl221	0/1	CrashLoopBackOff	3	2m18s	192.168.3.5	worker26	<none>	<none>
kube-system	calico-node-bnj6m	1/1	Running	0	2m18s	192.168.1.73	k8s-master	<none>	<none>
kube-system	calico-node-ns99h	1/1	Running	0	2m18s	192.168.1.89	worker6	<none>	<none>
kube-system	calico-node-grtbh	1/1	Running	0	2m18s	192.168.1.75	worker26	<none>	<none>
kube-system	coredns-74ff55c5b-bt6lb	0/1	Running	0	36m	192.168.3.2	worker26	<none>	<none>
kube-system	coredns-74ff55c5b-xqkt7	0/1	Running	0	36m	192.168.2.2	worker6	<none>	<none>
kube-system	etcd-k8s-master	1/1	Running	0	37m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-apiserver-k8s-master	1/1	Running	0	37m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-controller-manager-k8s-master	1/1	Running	0	37m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-proxy-d2rcc	1/1	Running	0	30m	192.168.1.89	worker6	<none>	<none>
kube-system	kube-proxy-mkz62	1/1	Running	0	30m	192.168.1.75	worker26	<none>	<none>
kube-system	kube-proxy-vbwvk	1/1	Running	0	36m	192.168.1.73	k8s-master	<none>	<none>
kube-system	kube-scheduler-k8s-master	1/1	Running	0	37m	192.168.1.73	k8s-master	<none>	<none>

- Scale down replica

```
]# kubectl scale deployment net-test --replicas=6
deployment.apps/net-test scaled
```

```
[root@k8s-master ~]# kubectl get pods -A -o wide | grep net-test
```

default	net-test-dc8bf8666-2g4js	1/1	Running	0	21m	192.168.2.8	worker6	<none>	<none>
default	net-test-dc8bf8666-4plqc	1/1	Running	0	21m	192.168.2.5	worker6	<none>	<none>
default	net-test-dc8bf8666-9rbz4	1/1	Running	0	21m	192.168.2.6	worker6	<none>	<none>
default	net-test-dc8bf8666-bvbf	1/1	Running	0	21m	192.168.2.4	worker6	<none>	<none>
default	net-test-dc8bf8666-jcdtt	1/1	Running	0	21m	192.168.2.7	worker6	<none>	<none>
default	net-test-dc8bf8666-pmx49	1/1	Running	0	21m	192.168.2.3	worker6	<none>	<none>

```
[root@k8s-master ~]#
```

```
[root@k8s-master ~]# kubectl scale deployment net-test --replicas=3
deployment.apps/net-test scaled
```

```
[root@k8s-master ~]# kubectl get pods -A -o wide | grep net-test
```

default	net-test-dc8bf8666-4plqc	1/1	Running	0	21m	192.168.2.5	worker6	<none>	<none>
default	net-test-dc8bf8666-bvbf	1/1	Running	0	21m	192.168.2.4	worker6	<none>	<none>
default	net-test-dc8bf8666-pmx49	1/1	Running	0	21m	192.168.2.3	worker6	<none>	<none>

```
[root@k8s-master ~]#
```

- Scale up replicas again

```
[root@k8s-master ~]# kubectl scale deployment net-test --replicas=6
deployment.apps/net-test scaled
```

```
[root@k8s-master ~]# kubectl get pods -A -o wide | grep net-test
```

default	net-test-dc8bf8666-4gm5x	0/1	ContainerCreating	0	7s	<none>	worker26	<none>	<none>
default	net-test-dc8bf8666-4plqc	1/1	Running	0	23m	192.168.2.5	worker6	<none>	<none>
default	net-test-dc8bf8666-bvbf	1/1	Running	0	23m	192.168.2.4	worker6	<none>	<none>
default	net-test-dc8bf8666-f4tvb	0/1	ContainerCreating	0	7s	<none>	worker26	<none>	<none>
default	net-test-dc8bf8666-pmx49	1/1	Running	0	23m	192.168.2.3	worker6	<none>	<none>
default	net-test-dc8bf8666-th5hf	1/1	Running	0	7s	192.168.187.197	worker26	<none>	<none>

```
[root@k8s-master ~]#
```

- Bring replicas to 0 and then to 10

```
[root@k8s-master ~]# kubectl get pods -A -o wide | grep net-test
```

default	net-test-dc8bf8666-71727	1/1	Running	0	33s	192.168.187.206	worker26	<none>	<none>
default	net-test-dc8bf8666-cgmdx	1/1	Running	0	33s	192.168.187.204	worker26	<none>	<none>
default	net-test-dc8bf8666-dz7zw	1/1	Running	0	3m5s	192.168.187.202	worker26	<none>	<none>
default	net-test-dc8bf8666-gr4d6	1/1	Running	0	3m5s	192.168.187.200	worker26	<none>	<none>
default	net-test-dc8bf8666-k6q42	1/1	Running	0	33s	192.168.35.1	worker6	<none>	<none>
default	net-test-dc8bf8666-mg5hz	1/1	Running	0	3m5s	192.168.187.201	worker26	<none>	<none>
default	net-test-dc8bf8666-pgmwg	1/1	Running	0	3m5s	192.168.187.203	worker26	<none>	<none>
default	net-test-dc8bf8666-tkx7m	1/1	Running	0	33s	192.168.187.207	worker26	<none>	<none>
default	net-test-dc8bf8666-w4kx1	1/1	Running	0	33s	192.168.35.2	worker6	<none>	<none>
default	net-test-dc8bf8666-zclwv	1/1	Running	0	33s	192.168.187.205	worker26	<none>	<none>

```
[root@k8s-master ~]#
```

- Scale up to 15

```
[root@k8s-master ~]# kubectl get pods -A -o wide | grep net-test
```

default	net-test-dc8bf8666-71727	1/1	Running	0	100s	192.168.187.206	worker26	<none>	<none>
default	net-test-dc8bf8666-7zrsr	1/1	Running	0	20s	192.168.35.4	worker6	<none>	<none>
default	net-test-dc8bf8666-cgmdx	1/1	Running	0	100s	192.168.187.204	worker26	<none>	<none>
default	net-test-dc8bf8666-dz7zw	1/1	Running	0	4m12s	192.168.187.202	worker26	<none>	<none>
default	net-test-dc8bf8666-gr4d6	1/1	Running	0	4m12s	192.168.187.200	worker26	<none>	<none>
default	net-test-dc8bf8666-hg6sf	1/1	Running	0	20s	192.168.187.210	worker26	<none>	<none>
default	net-test-dc8bf8666-k6q42	1/1	Running	0	100s	192.168.35.1	worker6	<none>	<none>
default	net-test-dc8bf8666-mg5hz	1/1	Running	0	4m12s	192.168.187.201	worker26	<none>	<none>
default	net-test-dc8bf8666-mk6pw	1/1	Running	0	20s	192.168.187.208	worker26	<none>	<none>
default	net-test-dc8bf8666-pgmwg	1/1	Running	0	4m12s	192.168.187.203	worker26	<none>	<none>
default	net-test-dc8bf8666-qvzqr	1/1	Running	0	20s	192.168.35.3	worker6	<none>	<none>
default	net-test-dc8bf8666-tkx7m	1/1	Running	0	100s	192.168.187.207	worker26	<none>	<none>
default	net-test-dc8bf8666-v4hk5	1/1	Running	0	20s	192.168.187.209	worker26	<none>	<none>
default	net-test-dc8bf8666-w4kx1	1/1	Running	0	100s	192.168.35.2	worker6	<none>	<none>
default	net-test-dc8bf8666-zclwv	1/1	Running	0	100s	192.168.187.205	worker26	<none>	<none>

```
[root@k8s-master ~]#
```

- Scaled down the deployment to 4

- And deployed one more deployment in different namespace

```
[root@k8s-master ~]# kubectl create ns my-net
namespace/my-net created
```

```
[root@k8s-master ~]# kubectl get ns
```

NAME	STATUS	AGE
default	Active	103m
kube-node-lease	Active	103m
kube-public	Active	103m
kube-system	Active	103m
my-net	Active	2s

```
[root@k8s-master ~]#
```

- Scale up/down both the replicas so we get pods on both the workers


```
[root@k8s-master ~]# kubectl get pods -o wide | grep net
net-test-dc8bf8666-2xs26 1/1 Running 0 62s 192.168.35.14 worker6 <none> <none>
net-test-dc8bf8666-7zrsr 1/1 Running 0 63m 192.168.35.4 worker6 <none> <none>
net-test-dc8bf8666-k6q42 1/1 Running 0 64m 192.168.35.1 worker6 <none> <none>
net-test-dc8bf8666-n56pl 1/1 Running 0 62s 192.168.187.252 worker26 <none> <none>
net-test-dc8bf8666-phlhg 1/1 Running 0 62s 192.168.35.13 worker6 <none> <none>
net-test-dc8bf8666-q8kzk 1/1 Running 0 62s 192.168.187.253 worker26 <none> <none>
net-test-dc8bf8666-qvzqr 1/1 Running 0 63m 192.168.35.3 worker6 <none> <none>
net-test-dc8bf8666-r8kbj 1/1 Running 0 62s 192.168.35.12 worker6 <none> <none>
net-test-dc8bf8666-w4kxl 1/1 Running 0 64m 192.168.35.2 worker6 <none> <none>
[root@k8s-master ~]# kubectl get pods -n my-net -o wide | grep net
my-net-c9dc49cdd-7mgkw 1/1 Running 0 56s 192.168.35.15 worker6 <none> <none>
my-net-c9dc49cdd-8jk58 1/1 Running 0 56s 192.168.187.195 worker26 <none> <none>
my-net-c9dc49cdd-9dlwz 1/1 Running 0 56s 192.168.35.16 worker6 <none> <none>
my-net-c9dc49cdd-b87c1 1/1 Running 0 6m38s 192.168.187.211 worker26 <none> <none>
my-net-c9dc49cdd-ct4pk 1/1 Running 0 56s 192.168.35.19 worker6 <none> <none>
my-net-c9dc49cdd-dfgv5 1/1 Running 0 6m38s 192.168.187.212 worker26 <none> <none>
my-net-c9dc49cdd-fswbb 1/1 Running 0 56s 192.168.35.17 worker6 <none> <none>
my-net-c9dc49cdd-ggxzr 1/1 Running 0 56s 192.168.187.255 worker26 <none> <none>
my-net-c9dc49cdd-jsznh 1/1 Running 0 56s 192.168.187.194 worker26 <none> <none>
my-net-c9dc49cdd-k9fdw 1/1 Running 0 56s 192.168.35.18 worker6 <none> <none>
my-net-c9dc49cdd-kq4zt 1/1 Running 0 56s 192.168.187.193 worker26 <none> <none>
my-net-c9dc49cdd-nwwh8 1/1 Running 0 56s 192.168.187.254 worker26 <none> <none>
my-net-c9dc49cdd-zmfjz 1/1 Running 0 3m53s 192.168.35.11 worker6 <none> <none>
[root@k8s-master ~]#
```

Check inter pod communication across namespaces:

Pod from namespace default namespace are able to access service running on pod under namespace "my-net" namespace

```
[root@k8s-master ~]# kubectl get pods | grep net
```

```
net-test-dc8bf8666-2xs26 1/1 Running 0 51s
net-test-dc8bf8666-7zrsr 1/1 Running 0 63m
net-test-dc8bf8666-k6q42 1/1 Running 0 64m
net-test-dc8bf8666-n56pl 1/1 Running 0 51s
net-test-dc8bf8666-phlhg 1/1 Running 0 51s
net-test-dc8bf8666-q8kzk 1/1 Running 0 51s
net-test-dc8bf8666-qvzqr 1/1 Running 0 63m
net-test-dc8bf8666-r8kbj 1/1 Running 0 51s
net-test-dc8bf8666-w4kxl 1/1 Running 0 64m
[root@k8s-master ~]# kubectl get pods -o wide | grep net
net-test-dc8bf8666-2xs26 1/1 Running 0 62s 192.168.35.14 worker6
net-test-dc8bf8666-7zrsr 1/1 Running 0 63m 192.168.35.4 worker6
net-test-dc8bf8666-k6q42 1/1 Running 0 64m 192.168.35.1 worker6
net-test-dc8bf8666-n56pl 1/1 Running 0 62s 192.168.187.252 worker26
net-test-dc8bf8666-phlhg 1/1 Running 0 62s 192.168.35.13 worker6
net-test-dc8bf8666-q8kzk 1/1 Running 0 62s 192.168.187.253 worker26
net-test-dc8bf8666-qvzqr 1/1 Running 0 63m 192.168.35.3 worker6
net-test-dc8bf8666-r8kbj 1/1 Running 0 62s 192.168.35.12 worker6
net-test-dc8bf8666-w4kxl 1/1 Running 0 64m 192.168.35.2 worker6
[root@k8s-master ~]# kubectl get pods -n my-net -o wide | grep net
my-net-c9dc49cdd-7mgkw 1/1 Running 0 56s 192.168.35.15 worker6
my-net-c9dc49cdd-8jk58 1/1 Running 0 56s 192.168.187.195 worker26
my-net-c9dc49cdd-9dlwz 1/1 Running 0 56s 192.168.35.16 worker6
my-net-c9dc49cdd-b87c1 1/1 Running 0 6m38s 192.168.187.211 worker26
my-net-c9dc49cdd-ct4pk 1/1 Running 0 56s 192.168.35.19 worker6
my-net-c9dc49cdd-dfgv5 1/1 Running 0 6m38s 192.168.187.212 worker26
my-net-c9dc49cdd-fswbb 1/1 Running 0 56s 192.168.35.17 worker6
my-net-c9dc49cdd-ggxzr 1/1 Running 0 56s 192.168.187.255 worker26
my-net-c9dc49cdd-jsznh 1/1 Running 0 56s 192.168.187.194 worker26
my-net-c9dc49cdd-k9fdw 1/1 Running 0 56s 192.168.35.18 worker6
my-net-c9dc49cdd-kq4zt 1/1 Running 0 56s 192.168.187.193 worker26
my-net-c9dc49cdd-nwwh8 1/1 Running 0 56s 192.168.187.254 worker26
my-net-c9dc49cdd-zmfjz 1/1 Running 0 3m53s 192.168.35.11 worker6
[root@k8s-master ~]#
```

```
[root@worker26 ~]# logout
```

Connection to worker26 closed.

```
[root@k8s-master ~]# kubectl exec -it net-test-dc8bf8666-2xs26 /bin/bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future
ersion. Use kubectl exec [POD] -- [COMMAND] instead.
```

```
root@net-test-dc8bf8666-2xs26:/# hostname -i
```

```
192.168.35.14
```

```
root@net-test-dc8bf8666-2xs26:/# curl http://192.168.187.195
```

```
<!DOCTYPE html>
```

```
<html>
```

```
<head>
```

```
<title>Welcome to nginx!</title>
```

```
<style>
```

```
body {
```

```
width: 35em;
```

```
margin: 0 auto;
```

```
font-family: Tahoma, Verdana, Arial, sans-serif;
```

```
}
```

```
</style>
```

```
</head>
```

```
<body>
```

```
<h1>Welcome to nginx!</h1>
```

```
<p>If you see this page, the nginx web server is successfully installed
working. Further configuration is required.</p>
```

```
<p>For online documentation and support please refer to
```

```
<a href="http://nginx.org/">nginx.org</a>.<br>
```

```
Commercial support is available at
```

```
<a href="http://nginx.com/">nginx.com</a>.</p>
```

```
</body>
```

```
<p><em>Thank you for using nginx.</em></p>
```

```
</body>
```

```
</html>
```

```
root@net-test-dc8bf8666-2xs26:/#
```

Calico now manages the networks and IPs in my nodes and assigns a different network address for each node.

Network Policy

Install calicoctl

```
[root@k8s-master ~]# kubectl create -f networkpolicy.yaml
error: unable to recognize "networkpolicy.yaml": no matches for kind "NetworkPolicy" in version "projectcalico.org/v3"
[root@k8s-master ~]# curl -O -L https://github.com/projectcalico/calicoctl/releases/download/v3.17.1/calicoctl
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 633 100 633 0 0 805 0 --:--:-- --:--:-- --:--:-- 806
100 37.1M 100 37.1M 0 0 291k 0 0:02:10 0:02:10 --:--:-- 287k
[root@k8s-master ~]# chmod +x calicoctl
[root@k8s-master ~]#
```

copy command to /usr/local/bin

Scenario:

```
[root@k8s-master networkpolicy]# kubectl get pods -n my-net -o wide --show-labels
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES LABELS
blue-66f9b45584-kjxxx 1/1 Running 0 17m 192.168.187.203 worker26 <none> <none> app=blue,pod-template-hash=66f9b45584
blue-66f9b45584-pkgyv 1/1 Running 0 40m 192.168.187.209 worker26 <none> <none> app=blue,pod-template-hash=66f9b45584
my-net-c9dc49cdd-7mgkw 1/1 Running 0 5h10m 192.168.35.15 worker6 <none> <none> app=my-net,pod-template-hash=c9dc49cdd
my-net-c9dc49cdd-fswbb 1/1 Running 0 5h10m 192.168.35.17 worker6 <none> <none> app=my-net,pod-template-hash=c9dc49cdd
my-net-c9dc49cdd-zmfjz 1/1 Running 0 5h13m 192.168.35.11 worker6 <none> <none> app=my-net,pod-template-hash=c9dc49cdd
red-7748479b97-4jcs9 1/1 Running 0 3h8m 192.168.35.23 worker6 <none> <none> app=red,pod-template-hash=7748479b97
red-7748479b97-d6icb 1/1 Running 0 3h8m 192.168.35.24 worker6 <none> <none> app=red,pod-template-hash=7748479b97
[root@k8s-master networkpolicy]#
```

```
2 pods = label "app=blue"
3 pods = label "app=my-net"
2 pods = label "app=red"
```

```
** NS = my-net
```

1) Apply policy within NS

Ingress traffic to endpoints in the namespace: my-net with label app: red is allowed, only if it comes from a pod in the same namespace with app: blue, on port 80.

Allow traffic within default NS from pods:

- ☐ Incoming traffic coming to pods labeled "app: red"
- ☐ Incoming traffic coming from pods labeled "app: blue"

```
[root@k8s-master networkpolicy]# cat networkpolicy.yaml
apiVersion: projectcalico.org/v3
kind: NetworkPolicy
metadata:
  name: allow-http-80-from-blue
  namespace: my-net
spec:
  selector: app == 'red'
  ingress:
  - action: Allow
    protocol: TCP
    sources:
      selector: app == 'blue'
    destination:
      ports:
        - 80
[root@k8s-master networkpolicy]#
```

```
[root@k8s-master ~]# calicoctl create -f networkpolicy.yaml
Successfully created 1 'NetworkPolicy' resource(s)
[root@k8s-master ~]#
```

```
[root@k8s-master networkpolicy]# kubectl get pods -n my-net -l app=blue -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE             NOMINATED NODE   READINESS GATES
blue-66f9b45584-kxjxx              1/1     Running   0           3h9m  192.168.187.203 worker26         <none>           <none>
blue-66f9b45584-pkgbv              1/1     Running   0           54m   192.168.187.209 worker26         <none>           <none>
[root@k8s-master networkpolicy]# kubectl get pods -n my-net -l app=red -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE             NOMINATED NODE   READINESS GATES
red-7748479b97-4jcz9              1/1     Running   0           3h22m  192.168.35.23  worker6          <none>           <none>
red-7748479b97-d6tdb              1/1     Running   0           3h22m  192.168.35.24  worker6          <none>           <none>
[root@k8s-master networkpolicy]#
```

Results: within NS

- Source - app: blue -> destination - app: red
 - Traffic allowed

```
[root@k8s-master ~]# kubectl exec -it blue-66f9b45584-kxjxx -n my-net
error: you must specify at least one command for the container
[root@k8s-master ~]#
[root@k8s-master ~]# kubectl exec -it blue-66f9b45584-kxjxx -n my-net /bin/bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@blue-66f9b45584-kxjxx:/# curl http://192.168.35.23
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 350px;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
root@blue-66f9b45584-kxjxx:/#
```

- Once policy is applied, only specific traffic is allowed

- Source - app: my-net -> destination - app: red
 - Traffic NOT allowed

```
[root@k8s-master networkpolicy]# kubectl get pods -n my-net -l app=my-net -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE             NOMINATED NODE   READINESS GATES
my-net-c9dc49cdd-7mgkw              1/1     Running   0           5h30m  192.168.35.15  worker6          <none>           <none>
my-net-c9dc49cdd-fswbb              1/1     Running   0           5h30m  192.168.35.17  worker6          <none>           <none>
my-net-c9dc49cdd-zmfjz              1/1     Running   0           5h33m  192.168.35.11  worker6          <none>           <none>
[root@k8s-master networkpolicy]#
```

```
[root@k8s-master ~]# kubectl exec -it my-net-c9dc49cdd-7mgkw -n my-net /bin/bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@my-net-c9dc49cdd-7mgkw:/# curl http://192.168.35.23
curl: (7) Failed to connect to 192.168.35.23 port 80: Connection timed out
root@my-net-c9dc49cdd-7mgkw:/#
```

2) Apply policy across NS

*** Previous policy has blocked traffic across NS

Ingress traffic to endpoints in the namespace: my-net with label app: red is allowed, if it comes from a pod in the namespace with type: backend, on port 80.

- Allow traffic within default NS from pods:
 - Incoming traffic coming to pods labeled "app: red" in my-net NS
 - Incoming traffic coming from pods labeled "app: database" in others NS

```
[root@k8s-master networkpolicy]# cat ns-networkpolicy.yaml
apiVersion: projectcalico.org/v3
kind: NetworkPolicy
metadata:
  name: allow-http-80-from-my-net-ns
  namespace: my-net
spec:
  selector: app == 'red'
  ingress:
  - action: Allow
    protocol: TCP
    sources:
      selector: app == 'database'
      namespaceSelector: type == 'backend'
    destination:
      ports:
        - 80
[root@k8s-master networkpolicy]# calicoctl create -f ns-networkpolicy.yaml
Successfully created 1 'NetworkPolicy' resource(s)
[root@k8s-master networkpolicy]#
```

Scenario:

- 2 pod - app: database [NS - others]
- 2 pod - app: red [NS - my-net]

```
[root@k8s-master networkpolicy]# kubectl get ns --show-labels
NAME      STATUS   AGE    LABELS
default   Active   7h34m  <none>
kube-node-lease   Active   7h34m  <none>
kube-public       Active   7h34m  <none>
kube-system       Active   7h34m  <none>
my-net           Active   5h51m  type=frontend
othersn         Active   2m24s  type=backend
[root@k8s-master networkpolicy]#
```

```
[root@k8s-master networkpolicy]# kubectl get pods -n my-net -l app=red -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   READINESS GATES
red-7748479b97-4jcz9    1/1     Running   0          3h49m  192.168.35.23  worker6  <none>            <none>
red-7748479b97-d6tdb    1/1     Running   0          3h49m  192.168.35.24  worker6  <none>            <none>
[root@k8s-master networkpolicy]#
```

```
[root@k8s-master networkpolicy]# kubectl get pods -o wide -n othersn --show-labels
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   READINESS GATES   LABELS
database-f86b86bdc-bm9ck  1/1     Running   0          32s   192.168.187.214  worker26  <none>            <none>            app=database,pod-template-hash=f86b86bdc
database-f86b86bdc-bpj47  1/1     Running   0          32s   192.168.187.213  worker26  <none>            <none>            app=database,pod-template-hash=f86b86bdc
[root@k8s-master networkpolicy]#
```

Results: across NS

- Source - app: database [othersn] --> destination - app: red (my-net)
- Traffic IS ALLOWED allow NOW

```
[root@k8s-master ~]# kubectl exec -it database-f86b86bdc-bm9ck -n othersn /bin/bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@database-f86b86bdc-bm9ck:/# curl http://192.168.35.23
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
root@database-f86b86bdc-bm9ck:/#
```

3) Apply deny all EGRESS policy on NS

Scenario:

```
[root@k8s-master networkpolicy]# kubectl get pods -n outgoing --show-labels -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   READINESS GATES   LABELS
blue-7bb46df96d-95fkh    1/1     Running   0          25m   192.168.187.215  worker26  <none>            <none>            app=blue,pod-template-hash=7bb46df96d
[root@k8s-master networkpolicy]# kubectl get pods -n myns --show-labels -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   READINESS GATES   LABELS
myweb-855c667ff6-rf6qh    1/1     Running   0          35s   192.168.187.216  worker26  <none>            <none>            app=myweb,pod-template-hash=855c667ff6
[root@k8s-master networkpolicy]#
```

Now apply deny all egress policy on outgoing ns

```
[root@k8s-master networkpolicy]# cat denyall.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: outgoing
  namespace: outgoing
spec:
  podSelector: {}
  policyTypes:
  - Egress
[root@k8s-master networkpolicy]#
```

```
[root@k8s-master networkpolicy]# kubectl create -f denyall.yaml
networkpolicy.networking.k8s.io/outgoing created
[root@k8s-master networkpolicy]#
```

Before the policy

```
[root@k8s-master ~]# kubectl exec -it blue-7bb46df96d-95fkh -n outgoing /bin/bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@blue-7bb46df96d-95fkh:/# curl http://192.168.187.216
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
root@blue-7bb46df96d-95fkh:/#
```

After deny all policy:

```
root@blue-7bb46df96d-95fkh:/# curl http://192.168.187.216

curl: (7) Failed to connect to 192.168.187.216 port 80: Connection timed out
root@blue-7bb46df96d-95fkh:/#
root@blue-7bb46df96d-95fkh:/#
```

Ingress rule present?	Egress rule present?	Value
No	No	Ingress
Yes	No	Ingress
No	Yes	Egress
Yes	Yes	Ingress, Egress