

Node Mgmt

12 January 2021 11:43

Node Drain

- Can purposefully drain the node of all the workloads so that the workloads are moved to other nodes.
 - The node is also cordoned or marked as **unschedulable**.

```
$ kubectl drain worker1
```

- There is also another command called cordon. Cordon simply marks a node unschedulable. Unlike drain it does not terminate or move the pods on an existing node.

```
$ kubectl cordon worker1
```

=====

- When the node is back online after a maintenance, it is still unschedulable. You then need to uncordon it.

```
$ kubectl uncordon worker-1
```

2 worker nodes

- 1) Deploy 6 replicas (assumption is that we will have 3 replicas on each worker node)
- 2) \$ kubectl drain worker1
 - a. Checking the results

```
$ kubectl cordon worker1
```

Node Selector:

Firstly apply label to Node

```
[root@master RBAC]# kubectl get nodes
NAME        STATUS    ROLES    AGE   VERSION
master      Ready    control-plane,master   8d   v1.20.1
worker1     Ready    <none>    8d   v1.20.1
```

```
[root@master RBAC]# kubectl label nodes worker1 disktype=ssd
node/worker1 labeled
```

```
[root@master RBAC]# kubectl get nodes --show-labels
NAME        STATUS    ROLES    AGE   VERSION   LABELS
master      Ready    control-plane,master   8d   v1.20.1   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=master,kubernetes.io/os=linux,node-role.kubernetes.io/control-plane=node-role.kubernetes.io/master=
worker1     Ready    <none>    8d   v1.20.1   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,disktype=ssd,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker1,kubernetes.io/os=linux
```

Secondly define Node selector in pod config file

```
[root@master RBAC]# cat 01-nodeselector.yml
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    env: test
spec:
  containers:
    - name: nginx
      image: nginx
      imagePullPolicy: IfNotPresent
      nodeSelector:
        disktype: ssd
```

```
[root@master ~]# kubectl apply -f 01-nodeselector.yml
pod/nginx created
```

```
[root@master ~]# kubectl get pod nginx -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP        NODE    NOMINATED NODE   READINESS GATES
nginx   1/1     Running   0          13s   10.44.0.15 worker1   <none>           <none>
```

Limitation

- Cannot provide advance expressions
- We used a single label and selector to achieve our goal here. But what if our requirement is much more complex.
 - e.g. disktype not ssd or size like large, medium or not small etc.
- For this we have Node Affinity and Anti affinity

Node Affinity

- Conceptually similar to nodeSelector
 - Allows you to constrain which nodes your pod is eligible to be scheduled on, based on labels on the node

Two Types :

- requiredDuringSchedulingIgnoredDuringExecution**
- preferredDuringSchedulingIgnoredDuringExecution**

Node Affinity Types

Available:

requiredDuringScheduling**Ignored**DuringExecution**preferred**DuringScheduling**Ignored**DuringExecution

	DuringScheduling	DuringExecution
Type 1	Required	Ignored
Type 2	Preferred	Ignored

-- Labeled node as "size: large"

-- Pod operators:

- large
- medium

```
[root@master ~]# kubectl label nodes worker1 size=large
node/worker1 labeled
```

```
[root@master ~]# kubectl get nodes --show-labels
NAME        STATUS    ROLES    AGE   VERSION   LABELS
master      Ready     control-plane,master   8d    v1.20.1   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=master,kubernetes.io/os=linux,node-role.kubernetes.io/control-plane=node-role.kubernetes.io/master=worker1    Ready     <none>      8d    v1.20.1   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,disktype=ssd,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker1,kubernetes.io/os=linux,size=large
```

```
[root@master ~]# cat 01-NodeAffinity.yml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
  - name: data-processor
    image: nginx
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
        - matchExpressions:
          - key: size
            operator: In
            values:
            - large
            - medium
```

kubectl apply -f 01-NodeAffinity.yml

```
[root@master ~]# kubectl apply -f 01-NodeAffinity.yml
pod/myapp-pod created
```

kubectl get pods myapp-pod -o wide

```
[root@master ~]# kubectl get pods myapp-pod -o wide
NAME        READY   STATUS    RESTARTS   AGE   IP          NODE   NOMINATED NODE   READINESS GATES
myapp-pod   0/1     ContainerCreating   0       4s    <none>      worker1    <none>           <none>
```

```
[root@master ~]# cat 02-NodeAffinity.yml
apiVersion: v1
kind: Pod
metadata:
  name: small
spec:
  containers:
  - name: data-processor
    image: nginx
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
        - matchExpressions:
          - key: size
            operator: NotIn
            values:
            - large
            - medium
```

```
[root@master ~]# kubectl apply -f 02-NodeAffinity.yml
pod/small created
```

```
[root@master ~]# kubectl get pods small -o wide
NAME        READY   STATUS    RESTARTS   AGE   IP          NODE   NOMINATED NODE   READINESS GATES
small       0/1     Pending   0          11s   <none>      <none>    <none>           <none>
```

Taints (node) and Toleration (pod)

- Pod to node relationship and how you can restrict what pods are placed on what nodes.
- Taints and Tolerations are used to set restrictions on what pods can be scheduled on a node.
- Only pods which are tolerant to the particular taint on a node will get scheduled on that node.

3 Taint Effects

- **NoSchedule** - Kubernetes will not schedule the pod onto that node
- **PreferNoSchedule** - Kubernetes will try to not schedule the pod onto the node
- **NoExecute**: pods that do not tolerate the taint will be evicted immediately

- Key:
- Value:
- Effect: explained above (behavior)
- Operation: condition

Apply Taint

```
# kubectl taint nodes worker1 app=blue:NoSchedule
```

```
[root@master ~]# kubectl taint nodes worker1 app=blue:NoSchedule
node/worker1 tainted
```

Check Taint:

```
# kubectl describe node worker1 |grep Taint
```

```
[root@master ~]# kubectl describe node worker1 |grep Taint
Taints:              app=blue:NoSchedule
```

Tolerations are added to pods by adding a tolerations section in pod definition.

```
[root@master ~]# cat 01-Toleration.yml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
  - name: nginx-container
    image: nginx
  tolerations:
  - key: "app"
    operator: "Equal"
    value: "blue"
    effect: "NoSchedule"
```

```
[root@master ~]# kubectl apply -f 01-Toleration.yml
pod/myapp-pod created
```

```
[root@master ~]# kubectl get pods myapp-pod -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE      NOMINATED NODE   READINESS GATES
myapp-pod 1/1     Running   0           35s   10.44.0.10   worker1   <none>            <none>
```

```
[root@calicomaster node]# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE      NOMINATED NODE   READINESS GATES
myapp-pod 1/1     Running   0           13m   192.168.235.130   worker1   <none>            <none>
myapp-t   1/1     Running   0           3m14s   192.168.235.131   worker1   <none>            <none>
nginx     1/1     Running   0           20m   192.168.235.129   worker1   <none>            <none>
small     0/1     Pending   0           11m   <none>         <none>     <none>            <none>
[root@calicomaster node]# kubectl taint node worker1 app=blue:NoExecute
node/worker1 tainted
[root@calicomaster node]# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE      NOMINATED NODE   READINESS GATES
myapp-pod 0/1     Terminating   0           14m   <none>         worker1   <none>            <none>
myapp-t   0/1     Terminating   0           3m43s   <none>         worker1   <none>            <none>
nginx     0/1     Terminating   0           20m   192.168.235.129   worker1   <none>            <none>
small     0/1     Pending        0           12m   <none>         <none>     <none>            <none>
[root@calicomaster node]#
```

Now create pod without toleration,

```
# kubectl run nginx --image=nginx
```

```
[root@master ~]# kubectl run nginx --image=nginx
pod/nginx created
```

Not assigned to worker node and taint is applied on it.

```
# kubectl get pod nginx -o wide
```

```
[root@master ~]# kubectl get pod nginx -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP        NODE    NOMINATED NODE   READINESS GATES
nginx   0/1     Pending   0           20s   <none>    <none>   <none>            <none>
```

Now we are removing taint from node. After removing nginx pod will be assigned to worker node

```
# kubectl taint nodes worker1 app=blue:NoSchedule-
```

```
[root@master ~]# kubectl taint nodes worker1 app=blue:NoSchedule-  
node/worker1 untainted
```

Now Pod is assigned to worker node

```
# kubectl get pod nginx -o wide
```

```
[root@master ~]# kubectl get pod nginx -o wide  
NAME      READY   STATUS    RESTARTS   AGE   IP            NODE      NOMINATED NODE   READINESS GATES  
nginx     1/1     Running   0           5m51s  10.44.0.15    worker1   <none>           <none>
```