

Set up is running with one master and one worker node

```
[root@master ~]# kubectl get nodes -o wide
NAME      STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION        CONTAINER-RUNTIME
master    Ready     control-plane,master   78m   v1.20.1   192.168.148.107 <none>        CentOS Linux 7 (Core) 3.10.0-1160.11.1.el7.x86_64 docker://19.3.11
worker    Ready     worker    49m   v1.20.1   192.168.148.110 <none>        CentOS Linux 7 (Core) 3.10.0-1160.11.1.el7.x86_64 docker://19.3.11

[root@master ~]# kubectl get pods -A
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  coredns-74ff55c5b-bqzmd                1/1     Running   1           77m
kube-system  coredns-74ff55c5b-bwcqj                1/1     Running   1           77m
kube-system  etcd-master                             1/1     Running   1           77m
kube-system  kube-apiserver-master                   1/1     Running   1           77m
kube-system  kube-controller-manager-master          1/1     Running   1           77m
kube-system  kube-flannel-ds-8v7tz                   1/1     Running   1           49m
kube-system  kube-flannel-ds-lfmm6                   1/1     Running   1           77m
kube-system  kube-proxy-hbxl8                        1/1     Running   1           77m
kube-system  kube-proxy-x96t7                        1/1     Running   1           49m
kube-system  kube-scheduler-master                   1/1     Running   1           77m
[root@master ~]#
```

- Create the 2 namespaces for this lab we are taking dev and prod namespaces
 - Current list of namespaces
 - Create two names space yaml as shown

```
[root@master ~]# cat << EOF > dev_ns.yml
apiVersion: v1
kind: Namespace
metadata:
  name: dev
EOF
```

```
[root@master ~]# cat << EOF > prod_ns.yml
apiVersion: v1
kind: Namespace
metadata:
  name: prod
EOF
```

```
[root@master ~]# kubectl get ns
NAME                STATUS   AGE
default             Active   80m
kube-node-lease     Active   80m
kube-public          Active   80m
kube-system          Active   80m
[root@master ~]# cat dev_ns.yml
apiVersion: v1
kind: Namespace
metadata:
  name: dev
[root@master ~]# cat prod_ns.yml
apiVersion: v1
kind: Namespace
metadata:
  name: prod
[root@master ~]#
```

- Create two names space using the yaml

```
[root@master ~]# kubectl create -f dev_ns.yml
namespace/dev created
[root@master ~]# kubectl create -f prod_ns.yml
namespace/prod created
[root@master ~]# kubectl get ns
NAME                STATUS   AGE
default             Active   82m
dev                 Active   11s
kube-node-lease     Active   82m
kube-public          Active   82m
kube-system          Active   82m
prod                 Active   6s
[root@master ~]#
```

- Now create the deployment yaml each in both the namespaces created above

```
[root@master ~]# cat << EOF > yaml_files/app1-dev.yml
> apiVersion: apps/v1
> kind: Deployment
> metadata:
>   name: hello
>   namespace: dev
> spec:
>   replicas: 1
>   selector:
>     matchLabels:
>       app: hello
>   template:
>     metadata:
```

```

> labels:
>   app: hello
> spec:
>   containers:
>   - name: hello-world
>     image: paulbouwer/hello-kubernetes:1.5
>     ports:
>     - containerPort: 8080
> EOF
[root@master ~]# cat << EOF > yaml_files/app2-prod.yml
> apiVersion: apps/v1
> kind: Deployment
> metadata:
>   name: hello
>   namespace: prod
> spec:
>   replicas: 1
>   selector:
>     matchLabels:
>     app: hello
>   template:
>     metadata:
>       labels:
>       app: hello
>     spec:
>       containers:
>       - name: hello-world
>         image: paulbouwer/hello-kubernetes:1.5
>         ports:
>         - containerPort: 8080
> EOF

```

- Once YAML files are ready apply the deployments

```

[root@master ~]# kubectl get pods -n dev
No resources found in dev namespace.
[root@master ~]# kubectl get pods -n prd
No resources found in prd namespace.
[root@master ~]# kubectl apply -f yaml_files/app1-dev.yml
deployment.apps/hello created
[root@master ~]# kubectl apply -f yaml_files/app2-prod.yml
deployment.apps/hello created
[root@master ~]# kubectl get pods -n dev
NAME                                READY   STATUS              RESTARTS   AGE
hello-565c6dd64b-mrmj4             0/1     ContainerCreating   0           9s
[root@master ~]# kubectl get pods -n prd
No resources found in prd namespace.
[root@master ~]# kubectl get pods -n prod
NAME                                READY   STATUS              RESTARTS   AGE
hello-565c6dd64b-vzsvw             0/1     ContainerCreating   0          13s
[root@master ~]#

```

- Wait till both the pods turns running

```

[root@master ~]# kubectl get pods -A | egrep 'dev|prod'
dev      hello-565c6dd64b-mrmj4            1/1     Running   0           62s
prod     hello-565c6dd64b-vzsvw            1/1     Running   0           58s

```

- Now define the service for both the dev and prod deployments

- Define the service for dev deployment

```

[root@master ~]# cat << EOF > yaml_files/app1-devsvc.yml
> apiVersion: v1
> kind: Service
> metadata:
>   name: hello
>   namespace: dev
> spec:
>   type: NodePort
>   ports:
>   - port: 80
>     targetPort: 8080
>     nodePort: 30080
>   selector:
>     app: hello
> EOF

```

- Create the service

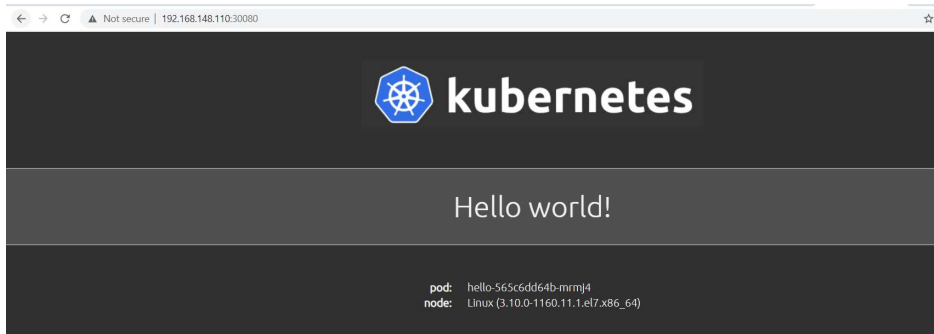
```

[root@master ~]# kubectl apply -f yaml_files/app1-devsvc.yml
service/hello created
[root@master ~]# kubectl get svc -n dev
NAME      TYPE        CLUSTER-IP      EXTERNAL-IP  PORT(S)          AGE
hello     NodePort    10.108.108.78    <none>       80:30080/TCP     7s
[root@master ~]#

```

- Now verify the application availability

```
[root@master ~]# kubectl describe svc -n dev
Name:          hello
Namespace:     dev
Labels:        <none>
Annotations:   <none>
Selector:      app=hello
Type:          NodePort
IP Families:   <none>
IP:            10.108.108.78
IPs:           10.108.108.78
Port:          <unset> 80/TCP
TargetPort:    8080/TCP
NodePort:      <unset> 30080/TCP
Endpoints:     192.168.1.6:8080
Session Affinity: None
External Traffic Policy: Cluster
Events:        <none>
```



```
[root@worker ~]# netstat -ntpl |grep 30080
tcp        0      0 0.0.0.0:30080          0.0.0.0:*               LISTEN     1665/kube-proxy
[root@worker ~]# curl http://localhost:30080 |grep world
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
100    654    100    654    0     0  33285      0 --:--:-- --:--:-- --:--:-- 34421
Hello world!
[root@worker ~]# curl http://192.168.148.110:30080 |grep world
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
100    654    100    654    0     0  24078      0 --:--:-- --:--:-- --:--:-- 25153
Hello world!
[root@worker ~]#
```

- Now repeat the same process for Prod deployment and use different port (we will use 30081)

```
[root@master ~]# cat << EOF > yam1_files/app1-prodsvc.yml
> apiVersion: v1
> kind: Service
> metadata:
>   name: hello
> namespace: prod
> spec:
>   type: NodePort
>   ports:
>     - port: 80
>       targetPort: 8080
>       nodePort: 30081
>   selector:
>     app: hello
> EOF
```

```
[root@master ~]# kubectl apply -f yam1_files/app1-prodsvc.yml
service/hello created
[root@master ~]# kubectl get svc -n prod
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
hello     NodePort  10.100.227.123   <none>           80:30081/TCP     15s
[root@master ~]# kubectl describe svc -n prod
Name:      hello
Namespace: prod
Labels:    <none>
Annotations: <none>
Selector:  app=hello
Type:      NodePort
IP Families: <none>
IP:        10.100.227.123
IPs:       10.100.227.123
Port:      <unset> 80/TCP
TargetPort: 8080/TCP
NodePort:   <unset> 30081/TCP
Endpoints:  192.168.1.7:8080
Session Affinity: None
External Traffic Policy: Cluster
Events:     <none>
```

```
[root@worker ~]# netstat -ntpl |grep 30081
tcp        0      0 0.0.0.0:30081          0.0.0.0:*               LISTEN     1665/kube-proxy
[root@worker ~]# curl http://192.168.148.110:30081 |grep world
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
100    654    100    654    0     0   7202      0 --:--:-- --:--:-- --:--:-- 7517
Hello world!
[root@worker ~]#
```

- Now create utility pod using busybox image which will be used for test the functionality

```
[root@master ~]# cat << EOF > yaml_files/utilitypod.yml
> apiVersion: v1
> kind: Pod
> metadata:
>   name: utility
> namespace: default
> spec:
>   containers:
>   - name: busybox
>     image: busybox:1.28
>     command:
>     - sleep
>     - "3600"
>   imagePullPolicy: IfNotPresent
>   restartPolicy: Always
> EOF
[root@master ~]#
```

- Apply the configuration and wait until pod gets in to the running state

```
[root@master ~]# kubectl apply -f yaml_files/utilitypod.yml
pod/utility created
[root@master ~]# kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
utility   1/1     Running   0          7s
[root@master ~]#
```

- Capture the Name and IP-Addr details

```
[root@master ~]# kubectl describe pods -n dev |egrep 'Name:|IP:'
Name:      hello-565c6dd64b-mrmj4
IP:        192.168.1.6
           192.168.1.6
SecretName: default-token-ws6nv
[root@master ~]# kubectl describe pods -n prod |egrep 'Name:|IP:'
Name:      hello-565c6dd64b-vzsvw
IP:        192.168.1.7
           192.168.1.7
SecretName: default-token-lqrr2
[root@master ~]# kubectl describe pods |egrep 'Name:|IP:'
Name:      utility
IP:        192.168.1.8
           192.168.1.8
SecretName: default-token-m47cj
[root@master ~]#
```

- Now verify the DNS Names

```
[root@master ~]# kubectl exec -it utility -- nslookup 192.168.1.6
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:        192.168.1.6
Address 1: 192.168.1.6 192-168-1-6.hello.dev.svc.cluster.local
[root@master ~]# kubectl exec -it utility -- nslookup 192.168.1.7
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:        192.168.1.7
Address 1: 192.168.1.7 192-168-1-7.hello.prod.svc.cluster.local
[root@master ~]# kubectl exec -it utility -- nslookup hello.dev.svc.cluster.local
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:        hello.dev.svc.cluster.local
Address 1: 10.108.108.78 hello.dev.svc.cluster.local
[root@master ~]# kubectl exec -it utility -- nslookup hello.prod.svc.cluster.local
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:        hello.prod.svc.cluster.local
Address 1: 10.100.227.123 hello.prod.svc.cluster.local
[root@master ~]#
```

- Now also try to ping the pod IP Address

```
[root@master ~]# kubectl exec -it utility -- ping 192.168.1.7 -c 1
PING 192.168.1.7 (192.168.1.7): 56 data bytes
64 bytes from 192.168.1.7: seq=0 ttl=64 time=0.295 ms

--- 192.168.1.7 ping statistics ---
1 packets transmitted, 1 packets received, 0% packet loss
round-trip min/avg/max = 0.295/0.295/0.295 ms
[root@master ~]# kubectl exec -it utility -- ping 192.168.1.6 -c 1
PING 192.168.1.6 (192.168.1.6): 56 data bytes
64 bytes from 192.168.1.6: seq=0 ttl=64 time=0.308 ms

--- 192.168.1.6 ping statistics ---
1 packets transmitted, 1 packets received, 0% packet loss
round-trip min/avg/max = 0.308/0.308/0.308 ms
[root@master ~]#
```